

Krish Ojha

11krishojha08@gmail.com | +91 83197 43894 | Indore, India

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

Backend engineer with 1+ year of production experience building scalable distributed systems and high-throughput APIs in Rust and TypeScript. Architected a Rust + Axum RFQ service with circuit-breaker fade tracking and sub-100ms response times under concurrent load, structured event-driven workflows for transaction processing and state management across PostgreSQL-backed microservices, and deployed fault-tolerant data infrastructure handling production-scale historical backfills. Comfortable across Rust (Axum, Tokio, Tower), Node.js (NestJS, Express), Python data pipelines, and React + TypeScript front-end.

SKILLS

Languages: Rust, TypeScript, JavaScript, Python, Java, Solidity, SQL

Backend & APIs: Rust (Axum, Tokio, Tower, Hyper), Node.js, NestJS, Express.js, RESTful APIs, Microservices, Event-Driven Architecture, Distributed Systems, JWT Authentication, Circuit Breakers, Retry Logic, Caching, Connection Pooling

Databases & Infrastructure: PostgreSQL, MongoDB, MySQL, Redis, Schema Design, Query Optimization, Indexing, State Machine Design, Docker, Docker Compose

Frontend: React.js, Next.js, Redux Toolkit, TypeScript-in-React, HTML5, CSS3, Tailwind CSS

Testing, Tools & DevOps: Git, GitHub Actions, Postman, npm, Vite, Foundry, Hardhat, Tenderly

Blockchain & Web3: Solidity, EVM, Cross-Chain Messaging, Bridge Integration, DeFi Protocols, Permit2, REVM, Flashbots

PROFESSIONAL EXPERIENCE

System Analyst

Ancilar | Feb 2025 - Present | Indore, India

- Constructed backend microservices in Rust (Axum, Tokio, Tower) and Node.js / TypeScript (NestJS) handling distributed transaction-processing workflows, including a high-throughput RFQ service with circuit-breaker fade tracking and sub-100ms API response times under concurrent load.
- Engineered resumable, fault-tolerant data ingestion and historical backfill systems featuring progress checkpointing, classified retry handling (transient, semi-permanent, permanent), idempotent upserts, and concurrent-update prevention, delivering 212,000+ rows of historical transaction data to production tables across 20+ source integrations spanning EVM and non-EVM blockchain ecosystems.
- Modeled complex transactional state machines in PostgreSQL across User, RFQ, Quoter, and Settlement service boundaries; led contract-first timeline planning and inter-service communication patterns ensuring consistency, idempotency, and reliability under high load.
- Contributed across the full stack of a production monitoring platform (Python data pipeline, NestJS API services, React frontend, Python SDK layer), surfacing 5 independent production bugs through systematic unit-test coverage and rolling out 3 substantive production fixes within a single month.
- Reviewed pull requests across 8+ engineers spanning frontend, backend, data pipeline, and SDK code, providing structured line-level feedback; established documentation standards including README, Postman collection, and design-considerations docs.
- Owned core backend for UniTokenSwap, an internal UniswapX-compatible RFQ trading platform: a Rust + Axum RFQ service with soft and hard RFQ semantics and circuit-breaker fallback, integrated with a Solidity reactor stack.

KEY PROJECTS

JobHunt - Multi-User Job-Search Platform

Personal project | Tech Stack: Python, FastAPI, SQLAlchemy, PostgreSQL, Groq LLM | Live: jobhunt-8i1m.onrender.com

- Built and deployed an end-to-end multi-user job platform on a fetch-once-match-many model, aggregating 10+ sources (REST APIs, ATS boards, public Telegram channels) into a bounded, de-duplicated catalog.
- Engineered an event-driven recommendation engine that learns from user actions (impressions, clicks, applies) via an online preference model and a blended relevance score, with an LLM re-rank pass.
- Designed resume-to-JD ATS matching, multi-channel delivery (Telegram, email), and a self-healing scheduler resilient to free-tier cold starts.
- Modeled the catalog, events, and per-user preference tables in PostgreSQL with idempotent upserts and URL-level de-duplication, so a single shared job pool serves N users at the cost of one.

Liquidity Data ETL Pipeline - Real-Time Data Processing System

Personal project | Tech Stack: Python, pandas, numpy, Multithreading, RESTful APIs

- Engineered a production-grade 3-tier ETL pipeline (Ingestor, Transform, API) following separation-of-concerns principles, handling large datasets with efficient ingestion and transformation workflows.
- Resolved multithreading challenges through thread-safe architecture with proper synchronization, achieving 5x throughput improvement over sequential processing.
- Developed an event listener system continuously tracking data sources and updating API responses with sub-500ms latency for always-fresh data availability.

Uniswap V3 Position Management Backend

Personal project | Tech Stack: TypeScript, Node.js, Ethers.js, REST APIs

- Built a 5-endpoint backend service for on-chain position management, covering state retrieval (liquidity, tick range, accumulated fees, position value), aggregation queries, and historical position endpoints.
- Engineered parallel RPC batching with response caching (~70% response time reduction), async timeout safety wrappers, real position-ID extraction from contract event logs, and clean REST contracts with structured error handling, request validation, and a Postman collection.
- Delivered clean REST contracts with structured error handling, request validation, and a Postman collection for fast integration testing.

EDUCATION

Acropolis Institute of Technology and Research | B.Tech, Computer Science & Data Science | Nov 2021 - May 2025 | 76%